

---

# Cache Friendly Shuffles for Machine Learning

---

May 8, 2016

## 1 OVERVIEW

## 2 LEAST SQUARES

## 3 WORD EMBEDDINGS

### 3.1 INTRODUCTION

In the word embeddings problem, given context counts  $X_{w,w'}$  we want to find word vectors  $v_w \in \mathbb{R}^k$  that minimizes the loss:

$$\min_{v,C} \sum_{w,w'} X_{w,w'} (\log(X_{w,w'}) - \|v_w + v_{w'}\|^2 - C)^2$$

### 3.2 EXPERIMENT DETAILS

We ran our experiments on the Edison compute nodes which feature two twelve core 2.4 GHz processors. However, we used only up to twelve cores/threads to avoid effects of NUMA. Word vectors were length 100 double arrays.

We used the first  $10^9$  bytes of English Wikipedia from <http://matmahoney.net/dc/textdata> as corpus data. After running the text preprocessing script supplied by the link, we computed co-occurrence counts of pairs of distinct words to create the parameter dependence graph. This graph was then fed into gpmets, computing a min-k-cut partitioning to create a cache-friendly ordering of the datapoints. k was set such that each block of k datapoints would reference just enough word vectors to fit into the L1-cache.

Hogwild was then run on the permuted co-occurrence graph generated by gpmets, maintaining the same ordering throughout execution. Although we experimented with both data sharding and no-data sharding, only results from data sharding are presented. To test hogwild without a cache-friendly shuffle, we randomly shuffled the datapoints.

We also ran the experiments on subsets of the corpus, repeating the procedure on the first 10%, 25%, 50% and 75% of the corpus data. In the full corpus data, there were 200,000 word vectors, and 30,000,000 datapoints.

## 4 CONCLUSION

## 5 FUTURE